

DOCKER 101

Why Docker?

Developing apps today requires so much more than writing code. Multiple languages, frameworks, architectures, and discontinuous interfaces between tools for each lifecycle stage creates enormous complexity.

For instance, if an end-to-end application is considered for the variety of tech stack it comprises, the maintenance becomes too difficult as a single tool upgrade might result in series of upgrades to other tool dependencies. This becomes a never ending vicious cycle. That's where Docker comes to the rescue. With docker each tool can be containerised with its very own libraries and dependencies such that others remain unaffected.

Docker simplifies and accelerates your workflow, while giving developers the freedom to innovate with their choice of tools, application stacks, and deployment environments for each project.

Containers are a standardized unit of software that allows developers to isolate their app from its environment, solving the "it works on my machine" headache. For millions of developers today, Docker is the de facto standard to build and share containerized apps - from desktop, to the cloud.

Getting started with Docker...

Docker can be easily installed on MacOS or any Linux based system. Docker host runs on Linux kernel on top of which containers are run.

Whereas when windows are in scope, a Linux virtual machine needs to be installed on which ,docker host is mounted on which containers will run.

Introduction to Images and Containers




IMAGE:

A Docker image is a read-only template that contains a set of instructions for creating a container that can run on the Docker platform. It provides a convenient way to package up applications and preconfigured server environments, which you can use for your own private use or share publicly with other Docker users.

A Docker image is made up of a collection of files that bundle together all the essentials – such as installations, application code, and dependencies – required to configure a fully operational container environment.

You can create a Docker image by using one of two methods:

Interactive: By running a container from an existing Docker image, manually changing that container environment through a series of live steps, and saving the resulting state as a new image.

Dockerfile: By constructing a plain-text file, known as a Dockerfile, which provides the specifications for creating a Docker image.

CONTAINER:

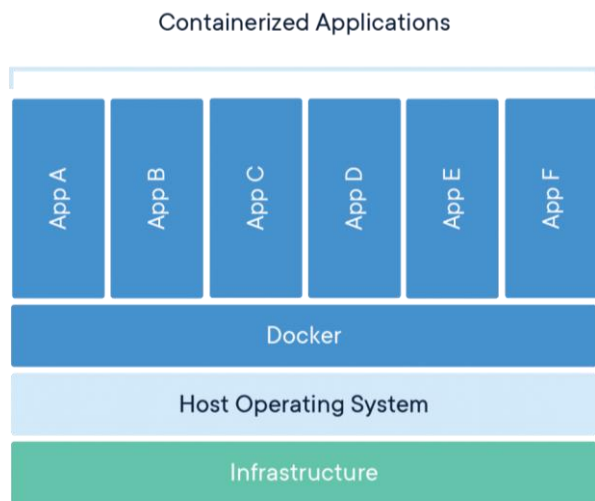
A Docker container is an open source software development platform. Its main benefit is to package applications in containers, allowing them to be portable to any system running a Linux or Windows operating system (OS). A container is created when a docker image is run on docker host. Multiple containers of the same image can be created with different port mapping.

***Please Note: The containers are supposed to host processes but not OS itself. So when the processes within a container complete, the container exits.

Containers vs Virtual Machines:

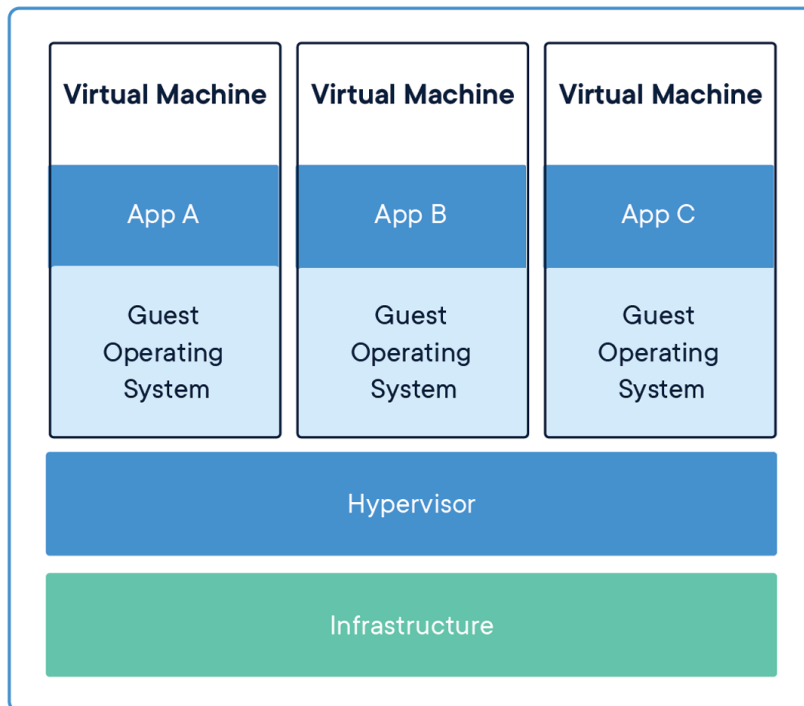
Containers:

Containers basically have the following structure:



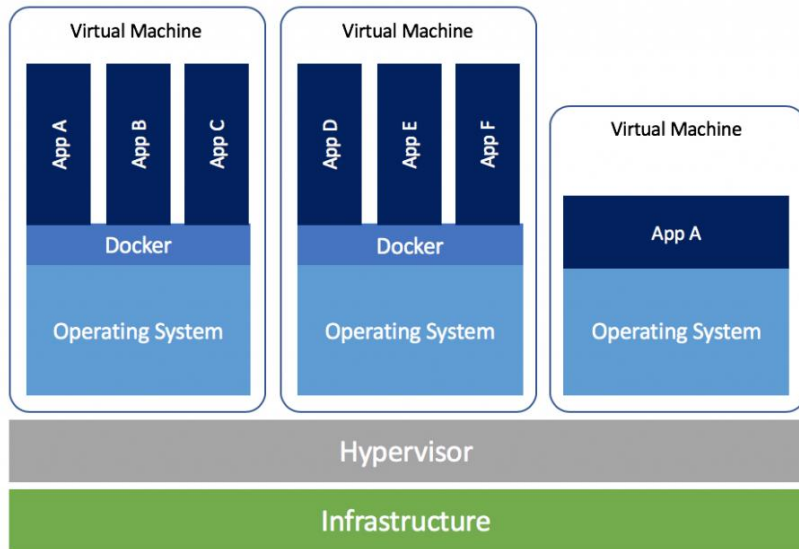
This structure allows less isolation. In other words, the containers share the same OS. This allows easy bootup. The memory too utilized will be comparatively less.

Virtual Machines:



Virtual Machines provide complete isolation to the applications. That is, each application has its own OS and each virtual machine is hosted on Hypervisor that interact with the hardware. This unfortunately puts in a lot of overhead and hence bootup takes time and the memory utilization too is expensive.

Best of both worlds:



In large environments, we can make use of both by integrating the best of both worlds as shown above.

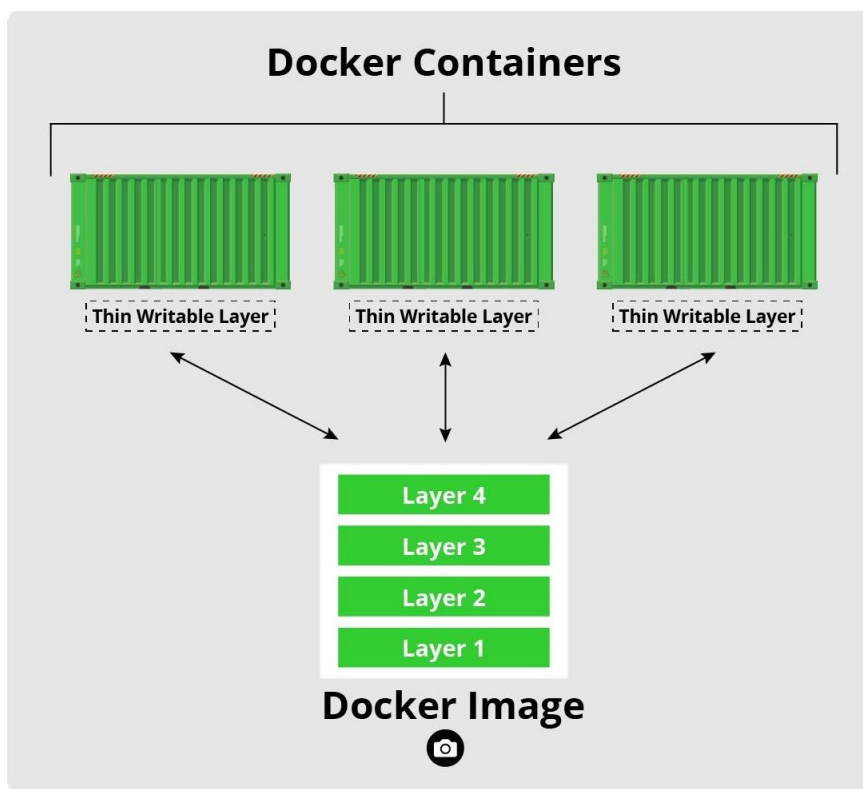
Anatomy of Images and Containers:

IMAGE LAYERS:

Each of the files that make up a Docker image is known as a layer. These layers form a series of intermediate images, built one on top of the other in stages, where each layer is dependent on the layer immediately below it. The hierarchy of your layers is key to efficient lifecycle management of your Docker images. Thus, you should organize layers that change most often as high up the stack as possible. This is because, when you make changes to a layer in your image, Docker not only rebuilds that particular layer, but all layers built from it. Therefore, a change to a layer at the top of a stack involves the least amount of computational work to rebuild the entire image.

CONTAINER LAYER:

Each time Docker launches a container from an image, it adds a thin writable layer, known as the container layer, which stores all changes to the container throughout its runtime. As this layer is the only difference between a live operational container and the source Docker image itself, any number of like-for-like containers can potentially share access to the same underlying image while maintaining their own individual state.



PARENT LAYER/BASE LAYER:

The first layer of a Docker image is known as the parent image. It's the foundation upon which all other layers are built and provides the basic building blocks for your container environments. Typically, parent layer can be an OS Version.

OR,

A base layer of a SQL image on which further layers of configurations can be built.

FILE SYSTEM DIFFERENCES BETWEEN IMAGES AND CONTAINERS

When you create a new container, you add a new writable layer on top of the underlying layers(read-only layers of the image). This layer is often called the "container layer". The major difference between a container and an image is this top writable layer. All changes made to the running container, such as writing new files, modifying existing files, and deleting files, are written to this thin writable container layer(diff).

When the container is deleted, the writable layer is also deleted. The underlying image remains unchanged. Because each container has its own writable container layer, and all changes are stored in this container layer, multiple containers can share access to the same underlying image and yet have their own data state.

DOCKER COMMANDS:

Docker run :

```
$ docker run [OPTIONS] IMAGE[:TAG|@DIGEST] [COMMAND] [ARG...]
```

1) DETACHED MODE:

When running a docker image , i.e., while creating container, you must first decide if you want to run the container in the background in a “detached” mode or in the default foreground mode.

which is -d = false

For instance:

```
$ docker run -d ubuntu
```

Ubuntu will run in detached mode.

2) FOREGROUND MODE:

- Interactive mode : -i , attaching STDIN
- Terminal mode : -t, attaching terminal
- Foreground mode: -a, adding STDIN,STDOUT,STDERR

For Instance:

Ubuntu container runs on open terminal as default which makes it to terminate it when the container processes are still running.

Hence,

Command will be like :

```
$ docker run -d ubuntu
```

This will create a container which has a 64 character long SHA-256 Id.

i.e., f78375b1c487e03c9438c729345e54db9d20cfa2ac1fc3494b6eb60872e74778

So to identify this container, we can either use the entire string of id or use the first few characters of it or the name of the container

i.e., f78 or f78375b1c487 or “evil_ptolemy”(name that’s generated randomly from docker host during container creation)

Suppose now its required to see the output of the container processes, we can re-attach the container foreground.

i.e.,

```
$ docker run -a f78
```

Or

```
$ docker run -a evil_ptolem
```

In case if we want to interact with the application, we can have an interactive mode along with terminal created:

Like:

```
$ docker run -it ubuntu bash
$ cat /etc/os-release
```

3) CICD SETUP:

Container ID can be passed to a specific file required as follows:

```
$ docker run -d ubuntu --cidfile /tmp/hello-world.cid
```

4) VERSION TAGS OF IMAGE:

A particular version of an image can be pulled from docker hub by specifying the tags

```
$ docker run -d ubuntu:22.04
```

5) NETWORK SETTING OF CONTAINER:

The network of a container can be set up with the following options.

None : the container will be isolated completely.

```
$ docker run --rm -d --network host --name my_nginx nginx
```

Bridge: Connect the container to the bridge. We can have multiple web server containers with bridge network nested on the same docker host, exposed with different port.

```
$ docker network create -d bridge my-bridge
```

```
$ docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
326ddef352c5	bridge	bridge	local
1ca18e6b4867	host	host	local
e0fc5f7ff50e	my-bridge	bridge	local
e9530f1fb046	none	null	local

```
$ docker run -d --name test1 --network my-bridge busybox sh -c "while true;do sleep 3600;done"
```

Host: uses network of that of the host. Hence we can have only single container of a webserver nested on the docker host in this case.

Compared to the default bridge mode, the host mode gives significantly better networking performance since it uses the host's native networking stack whereas the bridge has to go through one level of virtualization through the docker daemon. It is recommended to run containers in this mode when their networking performance is critical.

```
$ docker run -d --name test3 --net=host centos:7 /bin/bash -c "while true; do sleep 3600; done"
```

- Container: users network of another container specified.

```
$ docker run -d --name redis example/redis --bind 127.0.0.1
```

```
$ docker run --rm -it --network container:redis example/redis-cli -h 127.0.0.1
```

- User Defined Network:

```
$ docker network create -d bridge my-net
```

```
$ docker run --network=my-net -itd --name=container3 busybox
```

6) MANAGING /etc/hosts

Your container will have lines in /etc/hosts which define the hostname of the container itself as well as localhost and a few other common things. The --add-host flag can be used to add additional lines to /etc/hosts.

```
$ docker run -it --add-host db-static:86.75.30.9 ubuntu cat /etc/hosts
```

```
172.17.0.22    09d03f76bf2c
```

```
fe00::0       ip6-localnet
```



```
ff00::0    ip6-mcastprefix
ff02::1    ip6-allnodes
ff02::2    ip6-allrouters
127.0.0.1  localhost
::1        localhost ip6-localhost ip6-loopback
86.75.30.9 db-static
```

7) REMOVE WHEN EXITS

The container will be removed when the container exits.

```
$ docker run -it --rm --pid=host myhtop
```

8) RESTARTS WHEN EXITS

The container get restarted when exited. This can be mentioned explicitly when the container is created.

By default: --restart = no

Options :

no: default

```
$ docker run --restart=no redis
```

on-failure: restarts when it's a non-zero exit status and after a count of 10.

```
$ docker run --restart=on-failure:10 redis
```

always: restarts indefinitely

```
$ docker run --restart=always redis
```

unless-stopped: restarts on daemon start up and always regardless of exit status unless explicitly stopped.

```
$ docker run --restart=unless-stopped redis
```

9) OVERRIDING DOCKERFILE IMAGE DEFAULTS

CMD:

Consider a Dockerfile which is as follows:

```
FROM ubuntu
MAINTAINER xyz
RUN apt-get update
ENTRYPOINT ["echo", "Hello"]
CMD ["World"]
```

So when the image is built,

```
$ sudo docker build . -t ourimage
```

and run with the command :

```
$ sudo docker run ourimage
```

it prints :

“Hello World”

This happens as there is no name input given to the docker run command. So ideally it executes according to CMD given in the docker file.

- ENTRYPOINT:

When the docker run command is run with the name parameter:

i.e.,

```
$ sudo docker run ourimage xyz
```

it prints :

“Hello XYZ”

This is executed as per the entrypoint command which expects a name parameter.

So CMD is overridden by ENTRYPOINT command.

Entrypoint command can be executed from the docker run command as well.

i.e.,

```
$ docker run -it ourimage --entrypoint="Hello,xyz!"
```

This will override both ENTRYPOINT and CMD in Dockerfile and will print just “Hello,xyz!”

- EXPOSE:

This includes a few options that can be specified on the docker run command

`--expose=[]`

If you EXPOSE a port, the service in the container is not accessible from outside Docker, but from inside other Docker containers. So this is good for inter-container communication.

```
$ docker run --expose 5000 --name webserver nginx
```

`-P`

Publishes all ports to the host interfaces.

`-p=[]`

Publishes port to the docker host so the container is accessible from host/container

```
$ docker run -p 8080:5000 --name webserver nginx
```

`--link=""`

Helps to link containers to one another.

If we take a classic voting end-to-end application. it might have 2 dbs, one redis container named redis and one postgres container called DB. So the voter api will be connecting to both dbs. So to link these 3 we will have a command like :

```
$ docker run -d --link redis:redis --link db:db --name voter voter-app
```

- ENV:

Docker Automatically sets some environment variables when creating a linux container.

Such as :

HOME:	Set based on the value of USER
HOSTNAME	The hostname associated with the container
PATH	Includes popular directories as :

```
usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
```

TERM xterm if the container is allocated a pseudo-TTY

In case of additional env variables, it can be explicitly placed in command like :

```
$ export today=Wednesday
```

```
$ docker run -e "deep=purple" -e today --rm alpine env
```

- HEALTHCHECK:

Healthcheck is basically used to check health of the containers in intervals.

Such as:

```
$ docker run -d --name db --health-cmd "curl --fail http://localhost:8091/pools || exit 1" --health-interval=5s --health-retries=3 xyz/xyz
```

Multiple options that are available for us to setup the healthcheck for a container are as follows:

--health-cmd	Command to run to check health
--health-interval	Time between running the check
--health-retries	Consecutive failures needed to report unhealthy
--health-timeout	Maximum time to allow one check to run
--health-start-period	Start period for the container to initialize before starting health-retries countdown
--no-healthcheck	Disable any container-specified HEALTHCHECK

- VOLUMES:

For instance if we consider running a Jenkins container, when stopped the data is lost and the settings done should be redone again.

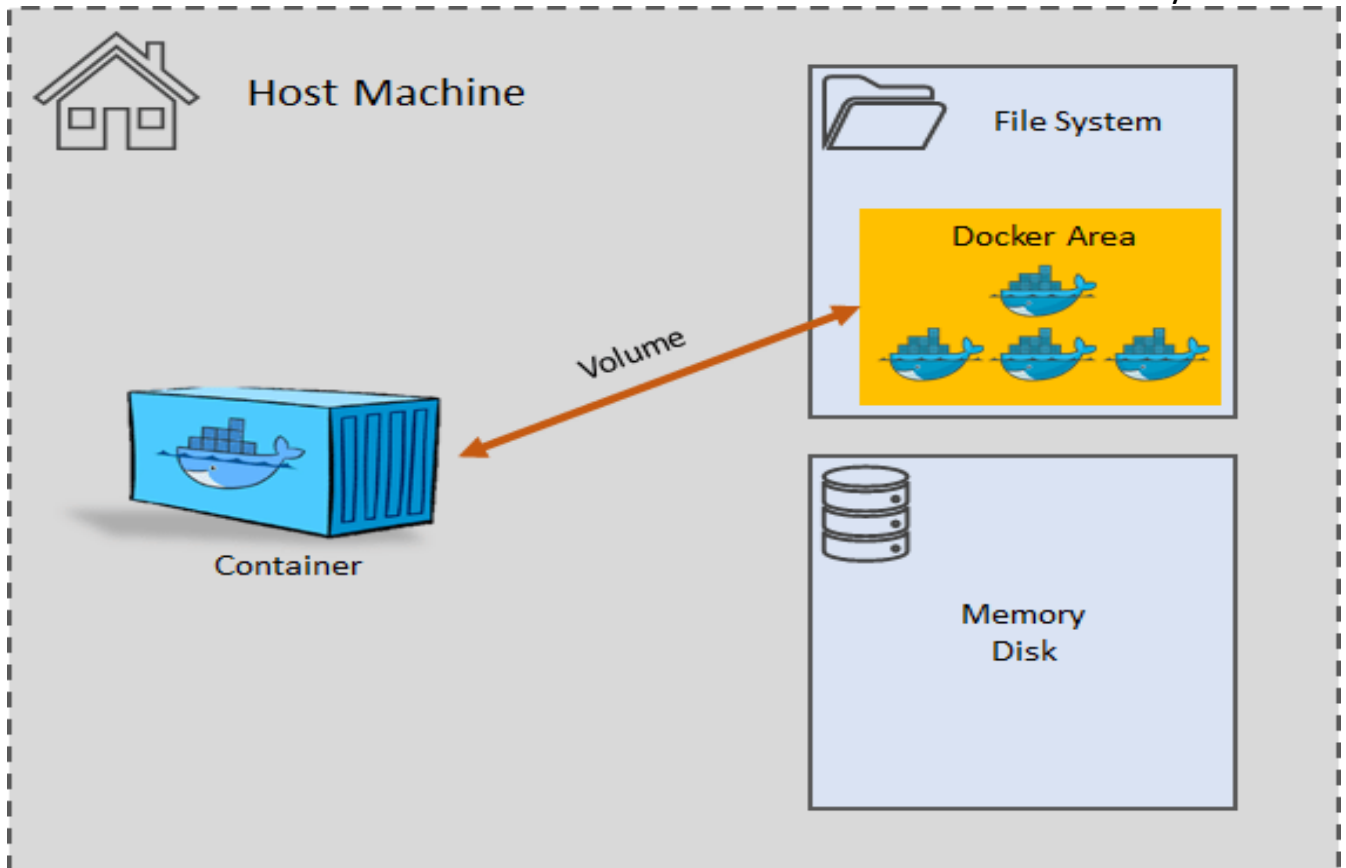
This is where volumes come into picture. We can store the data so that when it persists even when the container exits.

There are 3 types of volumes present

```
Docker Volumes
Bind Volumes
Tmpfs Mounts
```

Docker Volumes:

Volumes are part of the Host filesystem. For example, `/var/lib/jenkins` is the location that is shared with the docker container, and this location will act as the persistent data in the docker container. Remember this is isolated from the host machine to modify.



Scenarios that hold when using docker volumes:

- >> What happens when the container exits? We lose the data
- >> What happens when the docker daemon exits? We lose the data

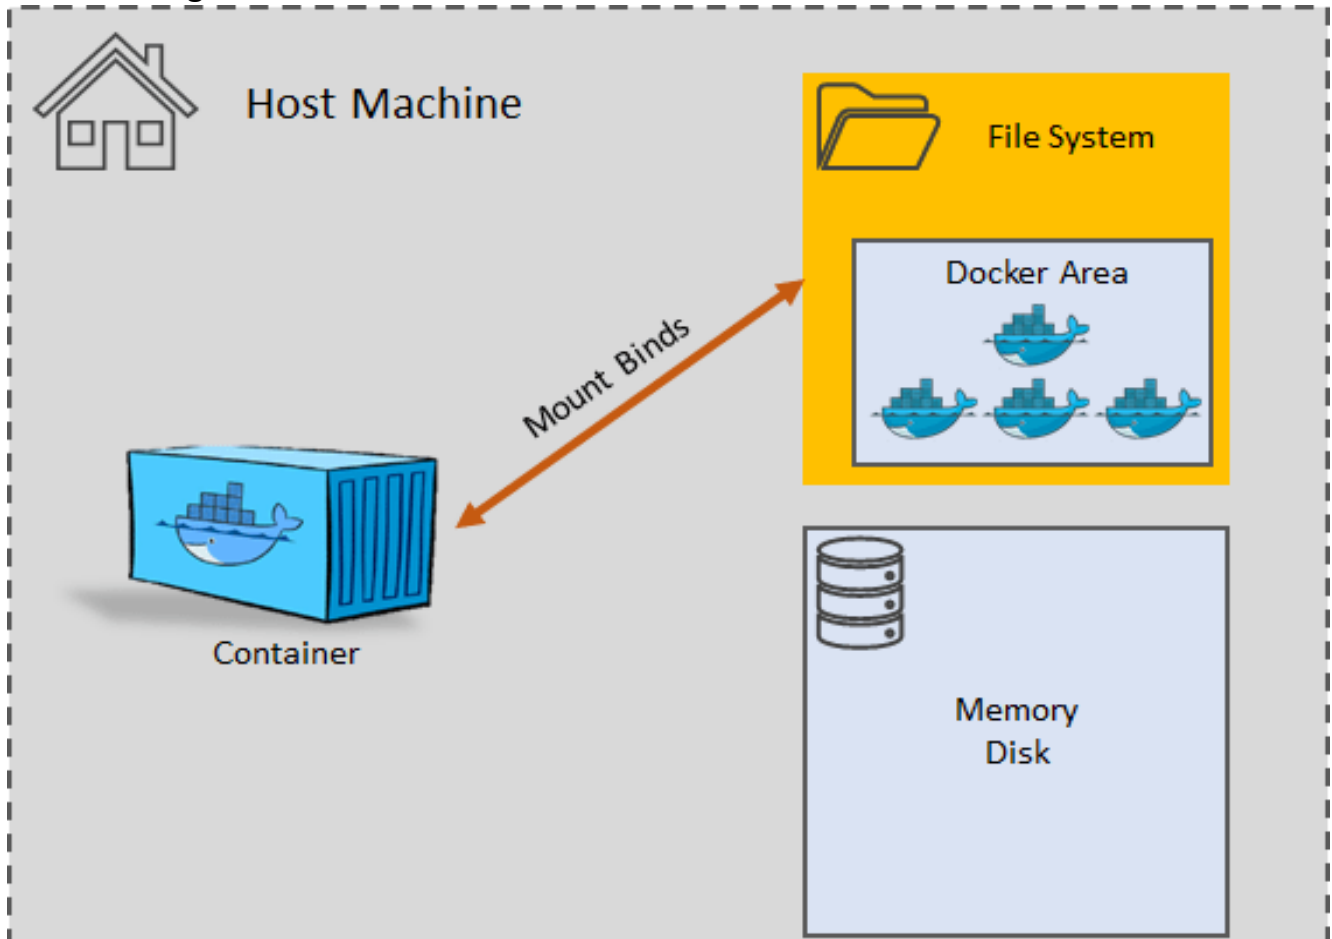
These Docker volumes can be created by command `docker volume create` command. Or this can be created when a container is created at any means like Dockerfile or Docker compose or docker command line.

Example:

```
$ docker run --name mysql --mount type=volume,src=volumename,dst=/container/path mysql:5.7
```

BIND MOUNTS:

Bind Mounts can be attached at any point at the file system or any disk which can already be used by the Host machine and can be modified from both containers and from the Host. But the limitation with the Bind Mounts are you cannot use the Docker CLI to manage Bind Mounts



Scenarios that hold when used Bind mounts:

>> What happens when the container exits? We don't lose the data

>> What happens when the docker daemon exits? We lose the data

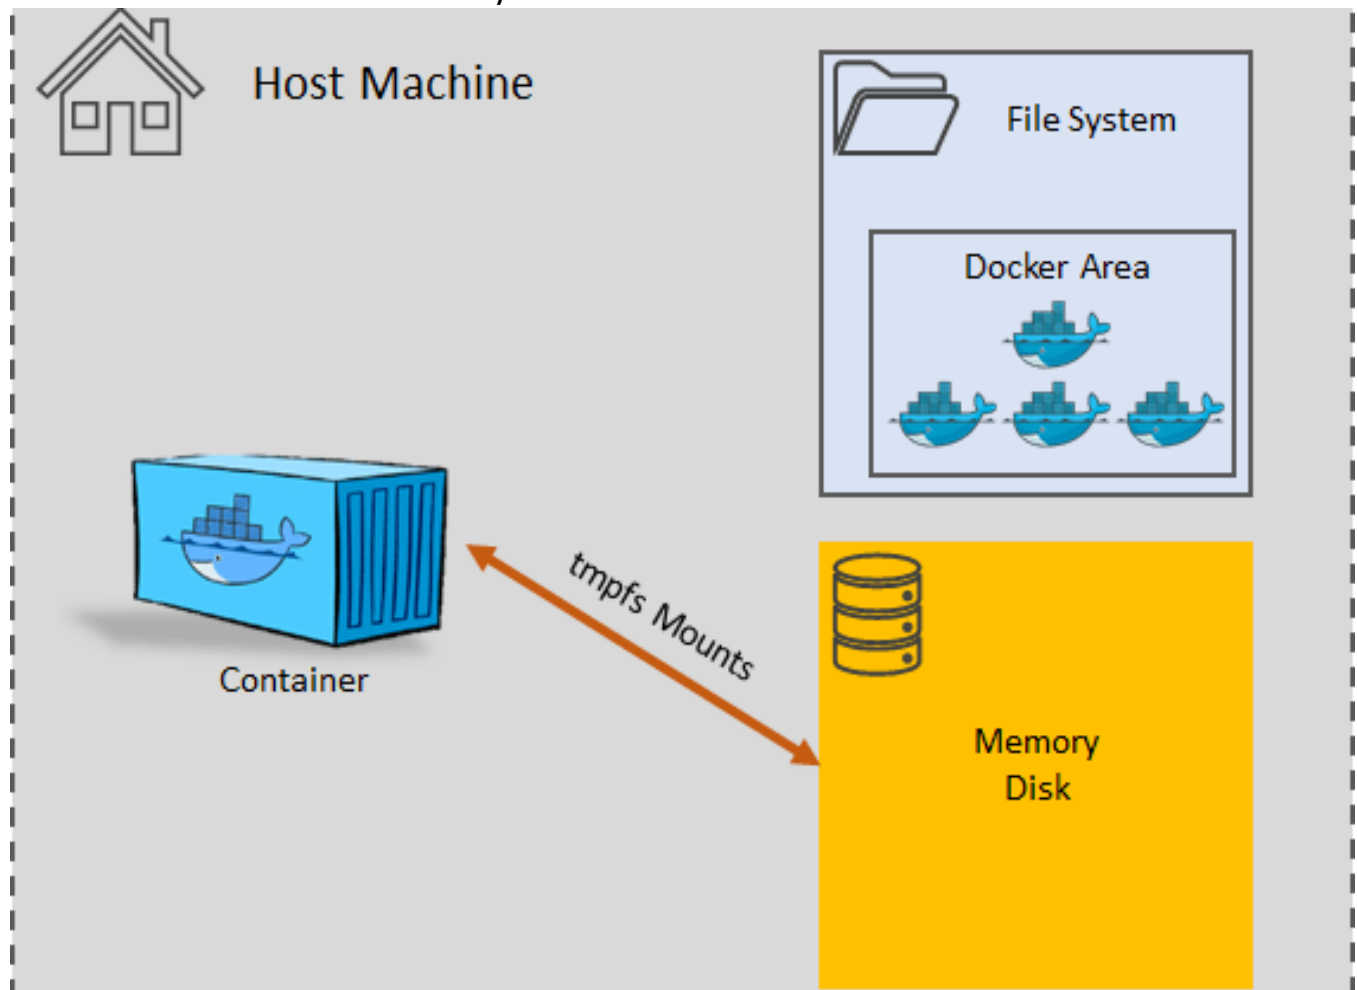
You can use `--mount type=bind, source=/host/path/,target=/container/path` to configure the Bind Mounts. Remember, Docker container can modify the file system attached with it and can be harmful if we misuse the containers

Example:

```
$ docker run --name mysql --mount type=bind,src=/host/path/,dst=/container/path mysql:5.7
```

TMPFS MOUNTS:

TMPFS mounts is not the persistent data on the disk. It will be available persistent on Neither Host nor Container filesystem.



Scenario that will hold when used TMPFS mounts:

>> What happens when the container exits? We don't lose the data

>> What happens when the docker daemon exits? We don't lose the data

You can create the mount by using `--mount type=tmpfs, destination=/app`. To create tmpfs mounts, you don't have to create a file structure in your host file system. You can mention only the destination file path and it will create the directory structure on its own.

Example:

```
$ docker run --name mysql --mount type=tmpfs, dst=/app/path mysql:5.7
```

Docker Attach:

Alternative for:

- `docker run -a`
- `docker run -it`

Example:

```
$ docker run -d --name topdemo ubuntu /usr/bin/top -b
$ docker attach topdemo
```

Docker Build:

This command is used to build the Dockerfile to create images of application that can be shipped and run anywhere.

1. Docker Build with URL:

```
$ docker build github.com/creack/docker-firefox
```

This will clone the git repository and use the cloned repository as context. The Dockerfile at the root will be used as “Dockerfile” to build the image. You can specify **git://** or **git@** scheme too.

```
$ docker build -f ctx/Dockerfile http://server/ctx.tar.gz
```

This sends the URL `http://server/ctx.tar.gz` to the Docker daemon, which downloads and extracts the referenced tarball. The `-f ctx/Dockerfile` parameter specifies a path inside `ctx.tar.gz` to the Dockerfile that is used to build the image

2. Docker Build with “-“:

```
docker build - < Dockerfile
```

This will read a Dockerfile from STDIN without context. Due to the lack of a context, no contents of any local directory will be sent to the Docker daemon. Since there is no context, a Dockerfile ADD only works if it refers to a remote URL.

```
$ docker build - < context.tar.gz
```

This will build an image for a compressed context read from STDIN. Supported formats are: bzip2, gzip and xz.

3. Docker Build with .dockerignorefile:

```
$ docker build .
```

docker build searches for a `.dockerignore` file relative to the Dockerfile name.

For example, running **docker build -f myapp.Dockerfile**. will first look for an ignore file named **myapp.Dockerfile.dockerignore**.

If such a file is not found, the `.dockerignore` file is used if present.

Using a Dockerfile based .dockerignore is useful if a project contains multiple Dockerfiles that expect to ignore different sets of files.

4. [Docker Build with tag “-t”:](#)

```
$ docker build -t vieux/apache:2.0 .
```

This will be building the image with the name provided with -t. We can provide multiple tag names given for an image that's built.

5. [Docker build with “-f”:](#)

```
$ docker build -f Dockerfile.colo .
```

This will use Dockerfile.colo to build the image instead of Dockerfile.

6. [Docker build with add-host](#)

This command is equivalent to running the image with bash and adding host into /etc/hosts

```
$ docker build --add-host=docker:10.180.0.1 .
```

7. [Docker build with --target](#)

The command allows us to specify an intermediate stage as final stage of the build of an image so that the rest of the stages can be ignored.

```
$ docker build -t mybuildimage --target build-env .
```

8. [Docker build with --output](#)

The command helps to specify the output such as a jar or tar file. The docker command usually builds an image. If we wish to build something else out of Dockerfile, we can do so with this option.

```
docker build --output type=tar,dest=out.tar .
```


Docker Commit:

This command is used to create an image out of a container.

```
$ docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
c3f279d17e0a	ubuntu:12.04	/bin/bash	7 days ago	Up 25 hours		desperate_dubinsky
197387f1b436	ubuntu:12.04	/bin/bash	7 days ago	Up 25 hours		focused_hamilton

```
$ docker commit c3f279d17e0a svendowideit/testimage:version3
```

```
f5283438590d
```

```
$ docker images
```

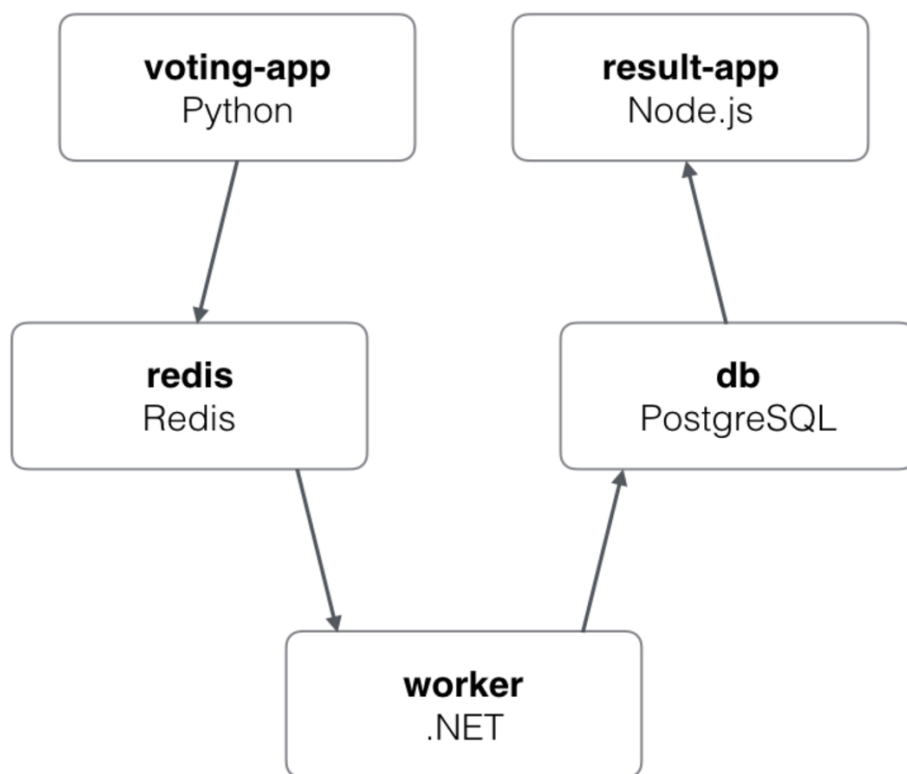
REPOSITORY	TAG	ID	CREATED	SIZE
svendowideit/testimage	version3	f5283438590d	16 seconds ago	335.7 MB

DOCKER COMPOSE

Docker compose yaml file is created when you want multiple containers to be brought up at a single time.

If we consider the classic voting application as example, it has multiple components.

- Vote – UI to cast vote
- Redis – in memory DB used as cache/message broker
- Vote worker – the backend that registers the vote and persists into the system.
- db – a classic PostgreSQL used to store the votes
- Result – UI to show the votes casted, to check who is leading/winning.



All these 5 elements are different images and to spun them up the classic way would require us to execute 5 different docker commands.

Instead, we can use docker compose yaml file to spin up all of these images together without hassle.

Docker commands to bring up this end to end application will be :

```
$ Docker run -d --name redis redis
```

```
$ Docker run -d --name db -e postgres_user postgres -e postgres_password password -p 5432:5432 postgres:9.4
```

```
$ Docker run -d --name vote -p 5000:80 --link redis:redis voting-app
```

```
$ Docker run -d --name result -p 5001:80 --link db:db result-app
```

```
$ Docker run -d --name worker --link db:db --link redis:redis worker
```

This will be a tedious process if we have more components which should be spun up on the account of an application.

To overcome this issue, we use docker compose.

So the same commands stated above can be decoded as below in docker compose and when executed, all of the components can be spun up implicitly.

DOCKER COMPOSE *(Please check the indentation, you might have to format it)*

```
version :2
  services
    redis
      image:redis
    db
      image:postgres:10.5
      port
        -5432:5432
      environment
        -postgres_user:postgres
        -postgres_password:password
    vote
      image:voter-app
      port
        -5000:80
      link
        -redis
    worker
      image:worker
      links
        -redis
        -db
    result
      image:result
      ports
        -5001:80
      links
        -db
```

Execution of this file is like :

```
$ docker compose up
```

When executed, all the containers will be created and we can have our application accessed at

- <http://localhost:5000> which will be to cast vote and
- <http://localhost:5001> which will show the results.

There are different versions of docker compose files available in which version 1 is deprecated.

The VERSION is usually specified at the top of the docker compose file. The latest version of docker compose for the same voting-app will comprise of additional bits of information.

The usual parts of docker compose has the following bits:

- VERSION
- SERVICES
- VOLUMES
- NETWORKS

(Please check the indentation, you might have to format it)

```
version: "3.9"
```

```
services:
```

```
  redis:
```

```
    image: redis:alpine
```

```
    ports:
```

```
      - "6379"
```

```
    networks:
```

```
      - frontend
```

```
  deploy:
```

```
    replicas: 2
```

```
    update_config:
```

```
      parallelism: 2
```

```
      delay: 10s
```

```
    restart_policy:
```

```
      condition: on-failure
```

```
  db:
```

```
    image: postgres:9.4
```

```
    volumes:
```

```
      - db-data:/var/lib/postgresql/data
```

```
    networks:
```

```
      - backend
```

deploy:

placement:

max_replicas_per_node: 1

constraints:

- "node.role==manager"

vote:

image: dockersamples/examplevotingapp_vote:before

ports:

- "5000:80"

networks:

- frontend

depends_on:

- redis

deploy:

replicas: 2

update_config:

parallelism: 2

restart_policy:

condition: on-failure

result:

image: dockersamples/examplevotingapp_result:before

ports:

- "5001:80"

networks:

- backend

depends_on:

- db

deploy:

replicas: 1

update_config:

parallelism: 2

delay: 10s

restart_policy:

condition: on-failure

worker:

image: dockersamples/examplevotingapp_worker

networks:

- frontend
- backend

deploy:

mode: replicated

replicas: 1

labels: [APP=VOTING]

restart_policy:

condition: on-failure

delay: 10s

max_attempts: 3

window: 120s

placement:

constraints:

- "node.role==manager"

visualizer:

image: dockersamples/visualizer:stable

ports:

- "8080:8080"

stop_grace_period: 1m30s

volumes:

- "/var/run/docker.sock:/var/run/docker.sock"

deploy:

placement:

constraints:

- "node.role==manager"

networks:

frontend:

backend:

volumes:

db-data:

BUILD:

- Build: Build can be specified either as a string containing a path to the build context.
- Context: Context is either a path to a directory containing a Dockerfile or a url to a Git repository
- Dockerfile: Dockerfile is the alternative file to build the image with.
- Args: Args will have all the build arguments required during the build process.
- Cache_from: Cache_from is used to the list of the images that engine uses for cache resolution.
- Labels: Add metadata to the resulting image using docker labels.
- Network: Set the network containers connect to the run instructions during the build.
- Shm_size: Set the size of the /dev/shm partition for the build containers.
- Target: Set the targeted stage. Such as dev, pre-prod or prod.
- Command: Command to override the default command.

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        - alpine: latest
        - corp/web_app: 3.14
      labels:
        com.example.description: "Accounting webapp"
        com.example.department: "Finance"
        com.example.label-with-empty-value: ""
    network: host
```

```
shm_size: '2gb'
target: prod
command: bundle exec thin -p 3000
```

CONFIG:

The top-level configs declaration defines or references configs that can be granted to the services in this stack. The source of the config is either file or external.

The configs have

- short syntax: this only specifies the name.
- long syntax: this specifies more granularity.

Thus the docker compose file would become :

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
    cache_from:
      - alpine: latest
      - corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      - source: my_config
        target: /redis_config
        uid: '103'
```



```
gid: '103'
mode: 0440
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

CONTAINER NAME

This is to specify a custom container name rather than a generated default random name. Hence Docker compose becomes:

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        - alpine: latest
        - corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      - source: my_config
        target: /redis_config
        uid: '103'
        gid: '103'
        mode: 0440
      container_name: custom_name
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
```

```
external: true
```

DEPENDS ON:

This option is nothing but `--link` in docker run command. It will help link the containers. So, possible docker compose will look like :

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
    cache_from:
      - alpine: latest
      - corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      - source: my_config
        target: /redis_config
        uid: '103'
        gid: '103'
        mode: 0440
    container_name: custom_name
    depends_on:
      - db
      - redis
```

```

redis:
  image: redis
db:
  image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true

```

DEPLOY:

Specifies how the container needs to be deployed.

- Endpoint_mode: can be either VIP or dns-rr(DNS-round robin)
- Labels: have labels specified for services.
- Mode: this can be either global/replicated.
 - If we have, say 3 components,/containers, for global, all the components will execute different tasks. Whereas, for replicas, each will be a replica of each other serving the same purpose.
- Placement: provides preferences, constraints and max number of replicas per container that it can spin up to.
 - constraints : such as
 - 1) The role of the user using or creating the container
 - 2) The type of the OS the container should run on
 - preferences:
 - spread: how the replicas should be spread across (zones)
 - max number of replicas
 - replicas: replicas that can be running at any given time.
 - Resources: provides the system limits that a container can have.
 - 1) CPUs
 - 2) memory
 - restart policy:
 - 1) condition: always/on-failure/none
 - 2) delay: how long to wait before restart
 - 3) max_attempts: how many times should the docker host try to bring up the container in case of failures.
 - 4) window: how long to wait before deciding if restart is succeeded or not.
 - Update_config:
 - 1) Parallelism: how many containers can be updated
 - 2) Delay: time between an update b/w a group of containers
 - 3) Failure_action: if we should continue, rollback or pause
 - 4) Monitor: how long should we monitor before the next task is rolled out

- 5) Max_failure_rate: failures to tolerate during an update
- 6) Order: how to carry the update
 1. stop-first: default; old task is stopped first before starting a new one.
 2. Start-first: start new task and then stop the old one

So the docker compose after adding deploy part will look as follows:

```
version: "3.9"

services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
    container_name: custom_name
    depends_on:
      db
      redis
    deploy:
      mode: replicated
      replicas: 6
      endpoint_mode: vip
```

```
  labels:
    com.example.description: "This label will appear on the web service"

  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04

  preferences:
    spread: node.labels.zone

  placement:
    max_replicas_per_node: 2

  update_config:
    parallelism: 2
    delay: 10s
    restart_policy:
      condition: on-failure

  resources:
    limits:
      cpus: '0.50'
      memory: 50M
    reservations:
      cpus: '0.25'
      memory: 20M
    restart_policy:
      condition: on-failure
      delay: 5s
      max_attempts: 3
      window: 120s
    update_config:
      parallelism: 2
      delay: 10s
    order: stop-first

  redis:
    image: redis

  db:
    image: postgres

configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

ENTRYPOINT:

Override the default entrypoint or command

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        alpine: latest
        corp/web_app: 3.14
      labels:
        com.example.description: "Accounting webapp"
        com.example.department: "Finance"
        com.example.label-with-empty-value: ""
      network: host
      shm_size: '2gb'
      target: prod
      command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
    container_name: custom_name
    depends_on:
      db
      redis
    deploy:
```

```
mode: replicated
replicas: 6
endpoint_mode: vip
labels:
com.example.description: "This label will appear on the web service"
constraints:
  node.role: manager
  engine.labels.operatingsystem: ubuntu 18.04
preferences:
  spread: node.labels.zone
placement:
  max_replicas_per_node: 2
update_config:
  parallelism: 2
  delay: 10s
restart_policy:
  condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
redis:
  image: redis
db:
  image: postgres
configs:
  my_config:
```

```
file: ./my_config.txt
my_other_config:
  external: true
```

ENV_FILE:

Adding a env variable file either in the same directory path or relative path.

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        alpine: latest
        corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
  configs:
    source: my_config
    target: /redis_config
    uid: '103'
    gid: '103'
    mode: 0440
  container_name: custom_name
  depends_on:
    db
    redis
  deploy:
```



```
mode: replicated
replicas: 6
endpoint_mode: vip
labels:
com.example.description: "This label will appear on the web service"
constraints:
  node.role: manager
  engine.labels.operatingsystem: ubuntu 18.04
preferences:
  spread: node.labels.zone
placement:
  max_replicas_per_node: 2
update_config:
  parallelism: 2
  delay: 10s
restart_policy:
  condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
redis:
  image: redis
db:
```

```
    image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

ENVIRONMENT:

Adding environment variables to the container in docker compose file

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        alpine: latest
        corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
    container_name: custom_name
    depends_on:
      db
      redis
    deploy:
```

```
mode: replicated
replicas: 6
endpoint_mode: vip
labels:
com.example.description: "This label will appear on the web service"
constraints:
  node.role: manager
  engine.labels.operatingsystem: ubuntu 18.04
preferences:
  spread: node.labels.zone
placement:
  max_replicas_per_node: 2
update_config:
  parallelism: 2
  delay: 10s
restart_policy:
  condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
environment:
  RACK_ENV: development
  SHOW: true
```

```
redis:
  image: redis
db:
  image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

EXPOSE:

This is used to expose the ports of a container to the host. However these ports will not be accessible outside the container.

This is equivalent to `-expose` in `docker run` command

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
    container_name: custom_name
    depends_on:
```

```
    db
    redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
  com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
```

```
environment:
  RACK_ENV: development
  SHOW: true
  expose:
    - "3000"
redis:
  image: redis
db:
  image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

EXTERNAL LINKS:

This is used to link with the containers outside the docker-compose.yml.

So the docker compose will become like :

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
```

```
    target: /redis_config
    uid: '103'
    gid: '103'
    mode: 0440
container_name: custom_name
depends_on:
    db
    redis
deploy:
    mode: replicated
    replicas: 6
    endpoint_mode: vip
    labels:
com.example.description: "This label will appear on the web service"
constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
preferences:
    spread: node.labels.zone
placement:
    max_replicas_per_node: 2
update_config:
    parallelism: 2
    delay: 10s
restart_policy:
    condition: on-failure
resources:
    limits:
        cpus: '0.50'
        memory: 50M
    reservations:
        cpus: '0.25'
        memory: 20M
restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
update_config:
    parallelism: 2
```

```
    delay: 10s
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
  - ./common.env
  - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  redis:
    image: redis
  db:
    image: postgres
  configs:
    my_config:
      file: ./my_config.txt
    my_other_config:
      external: true
```

HEALTHCHECK:

Configure the check to run to determine the containers health if they are running as expected or if the container has exited.

- Test: string or an arraylist .If arraylist, first item will be none/cmd/cmd-shell
- Interval
- Timeout
- Retries
- Start_period

So the docker compose will become like :

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
```



```
cache_from:
  alpine: latest
  corp/web_app: 3.14
labels:
  com.example.description: "Accounting webapp"
  com.example.department: "Finance"
  com.example.label-with-empty-value: ""
network: host
shm_size: '2gb'
target: prod
command: bundle exec thin -p 3000
configs:
  source: my_config
  target: /redis_config
  uid: '103'
  gid: '103'
  mode: 0440
container_name: custom_name
depends_on:
  db
  redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
    com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
```


```
limits:
  cpus: '0.50'
  memory: 50M
reservations:
  cpus: '0.25'
  memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
environment:
  RACK_ENV: development
  SHOW: true
expose:
- "3000"
external_links:
  -project_db:db_mysql
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost"]
  interval: 1m30s
  timeout: 10s
  retries: 3
  start_period: 40s
redis:
  image: redis
db:
  image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
```

```
external: true
```

INIT


Run init inside a container that forwards the signals and reaps processes.(equivalent to docker run).This is responsible to start the system, ssh, bash and other daemons.

Why init is required?



Imagine our container runs a webserver that runs a CGI script that's written in bash.This script calls grep.

Consider an instance where web server decides that this cgi script is taking too long and kills the script.But grep is not affected and it keeps running.



When grep finishes, it becomes a zombie.

The web server doesn't know about grep and it doesn't reap it. Hence grep stays in system.

We usually run third party apps/containers. You are running someone else's code, so you really need to be sure that those apps don't spawn processes in such a way that become/create zombie processes later. Hence we should run a proper init system to prevent problems.

So the docker compose will become as follows:

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
```

```
    alpine: latest
    corp/web_app: 3.14
  labels:
    com.example.description: "Accounting webapp"
    com.example.department: "Finance"
    com.example.label-with-empty-value: ""
  network: host
  shm_size: '2gb'
  target: prod
  command: bundle exec thin -p 3000
  configs:
    source: my_config
    target: /redis_config
    uid: '103'
    gid: '103'
    mode: 0440
  container_name: custom_name
  depends_on:
    db
    redis
  deploy:
    mode: replicated
    replicas: 6
    endpoint_mode: vip
    labels:
      com.example.description: "This label will appear on the web service"
    constraints:
      node.role: manager
      engine.labels.operatingsystem: ubuntu 18.04
    preferences:
      spread: node.labels.zone
    placement:
      max_replicas_per_node: 2
    update_config:
      parallelism: 2
      delay: 10s
    restart_policy:
      condition: on-failure
  resources:
    limits:
```

```
    cpus: '0.50'
    memory: 50M
    reservations:
      cpus: '0.25'
      memory: 20M
  restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
  update_config:
    parallelism: 2
    delay: 10s
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
    - ./common.env
    - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost"]
    interval: 1m30s
    timeout: 10s
    retries: 3
    start_period: 40s
  redis:
    image: redis
    init: true
  db:
    image: postgres
  configs:
    my_config:
      file: ./my_config.txt
    my_other_config:
```

```
external: true
```

LOGGING:

This gives an option to configure logging for the services.

- Driver : syslog / json-file / awslogs
- Options :
 - Syslog-address
 - Max-size
 - Max-file

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
```

```
container_name: custom_name
depends_on:
  db
  redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
    com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
```

```
- ./common.env
- ./apps/web.env
environment:
  RACK_ENV: development
  SHOW: true
expose:
  - "3000"
external_links:
  -project_db:db_mysql
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost"]
  interval: 1m30s
  timeout: 10s
  retries: 3
  start_period: 40s
logging:
  driver: "json-file"
  options:
    max-size: "200k"
    max-file: "10"
redis:
  image: redis
  init: true
db:
  image: postgres
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

Some of the other options can be like :

```
logging:
  driver: "awslogs"
  options:
    awslogs-region: us-east-1

logging:
  driver: "gelf"
  options:
    gelf-address: "udp://1.2.3.4:12201"
```



```
logging:
  driver: "gelf"
  options:
    gelf-address: "udp://1.2.3.4:12201"
logging:
  driver: "syslog"
  options:
    syslog-address: "udp://1.2.3.4:12201"
```

NETWORK MODE

The options to this command are

- Bridge
- Host
- None
- Container : `network_mode: "container:[container name/id]"`
- Service : `network_mode: "service:[service name]"`

This is same as docker run with `–network` command.

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
```

```
    uid: '103'
    gid: '103'
    mode: 0440
  container_name: custom_name
  depends_on:
    db
    redis
  deploy:
    mode: replicated
    replicas: 6
    endpoint_mode: vip
    labels:
      com.example.description: "This label will appear on the web service"
    constraints:
      node.role: manager
      engine.labels.operatingsystem: ubuntu 18.04
    preferences:
      spread: node.labels.zone
    placement:
      max_replicas_per_node: 2
    update_config:
      parallelism: 2
      delay: 10s
    restart_policy:
      condition: on-failure
  resources:
    limits:
      cpus: '0.50'
      memory: 50M
    reservations:
      cpus: '0.25'
      memory: 20M
  restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
  update_config:
    parallelism: 2
    delay: 10s
```

```
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
    - ./common.env
    - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost"]
    interval: 1m30s
    timeout: 10s
    retries: 3
    start_period: 40s
  logging:
    driver: "json-file"
    options:
      max-size: "200k"
      max-file: "10"
  network_mode: "service:redis"
  redis:
    image: redis
    init: true
  db:
    image: postgres
  configs:
    my_config:
      file: ./my_config.txt
    my_other_config:
      external: true
```

NETWORKS:

Networks is a section similar to that of services.

- Aliases:

The container can use the network name given in the networks section or can have a alias described with “aliases”

```
version: "3.9"

services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
  configs:
    source: my_config
    target: /redis_config
    uid: '103'
    gid: '103'
```

```
mode: 0440
container_name: custom_name
depends_on:
  db
  redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
    com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
```

```
env_file:
- ./common.env
- ./apps/web.env
environment:
  RACK_ENV: development
  SHOW: true
expose:
  - "3000"
external_links:
  -project_db:db_mysql
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost"]
  interval: 1m30s
  timeout: 10s
  retries: 3
  start_period: 40s
logging:
  driver: "json-file"
  options:
    max-size: "200k"
    max-file: "10"
```

```
networks: new
```

```
redis:
  image: redis
  init: true
networks: legacy
```

```
db:
  image: postgres
networks:
```

```
new:
```

```
aliases:
```

```
- database
```

```
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
```

```
networks:
```

```
new:
```

```
legacy:
```

ipv4_address,ipv6_address

Specify ip address for containers for this service when joining the network.

```
version: "3.9"

services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
  configs:
    source: my_config
    target: /redis_config
    uid: '103'
    gid: '103'
    mode: 0440
  container_name: custom_name
  depends_on:
    db
```

```
    redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
    com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
environment:
```



```
RACK_ENV: development
SHOW: true
expose:
  - "3000"
external_links:
  -project_db:db_mysql
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost"]
  interval: 1m30s
  timeout: 10s
  retries: 3
  start_period: 40s
logging:
  driver: "json-file"
  options:
    max-size: "200k"
    max-file: "10"
networks:
  new:
    ipv4_address: 172.16.238.10
    ipv6_address: 2001:3984:3989::10
redis:
  image: redis
  init: true
  networks:
    new:
      ipv4_address: 172.16.238.20
      ipv6_address: 2001:3984:3989::20
db:
  image: postgres
  networks:
    legacy:
      ipv4_address: 172.16.238.45
      ipv6_address: 2001:3984:3989::45
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
networks:
```

```

new:
  ipam:
    driver: default
    config:
      - subnet: "172.16.238.0/24"
      - subnet: "2001:3984:3989::/64"
  legacy:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.25/48"
        - subnet: "2001:3984:3989::65/128"

```

PORTS

- Short syntax: this has 3 options
 - Specify both ports – host: container
"8000:8000"
 - Specify container port
"3000-3005"
 - Specify host ip address to bind to AND both ports
"127.0.0.1:8001:8001"
"127.0.0.1:5000-5010:5000-5010"
"127.0.0.1::5000"
- Long syntax: This has 4 options
 - Target
 - Published
 - Protocol
 - Mode

```

version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
    cache_from:
      alpine: latest
      corp/web_app: 3.14
    labels:

```

```
    com.example.description: "Accounting webapp"
    com.example.department: "Finance"
    com.example.label-with-empty-value: ""
network: host
shm_size: '2gb'
target: prod
command: bundle exec thin -p 3000
configs:
    source: my_config
    target: /redis_config
    uid: '103'
    gid: '103'
    mode: 0440
container_name: custom_name
depends_on:
    db
    redis
deploy:
    mode: replicated
    replicas: 6
    endpoint_mode: vip
    labels:
        com.example.description: "This label will appear on the web service"
constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
preferences:
    spread: node.labels.zone
placement:
    max_replicas_per_node: 2
update_config:
    parallelism: 2
    delay: 10s
restart_policy:
    condition: on-failure
resources:
    limits:
        cpus: '0.50'
        memory: 50M
    reservations:
```

```
    cpus: '0.25'
    memory: 20M
  restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
  update_config:
    parallelism: 2
    delay: 10s
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
    - ./common.env
    - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost"]
    interval: 1m30s
    timeout: 10s
    retries: 3
    start_period: 40s
  logging:
    driver: "json-file"
    options:
      max-size: "200k"
      max-file: "10"
  networks:
    new:
      ipv4_address: 172.16.238.10
      ipv6_address: 2001:3984:3989::10
  ports:
    - "172.16.238.10:5000:80"
  redis:
```

```
  image: redis
  init: true
db:
  image: postgres
  networks:
    new:
      aliases:
        - database
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
networks:
  new:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.0/24"
        - subnet: "2001:3984:3989::/64"
  legacy:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.25/48"
        - subnet: "2001:3984:3989::65/128"
```

PROFILES

Profiles can be used to label the components of an end to end components such as backend, frontend, mid-tier

```
version: "3.9"
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
    args:
      buildno: 1
    cache_from:
      alpine: latest
```

```
    corp/web_app: 3.14
labels:
  com.example.description: "Accounting webapp"
  com.example.department: "Finance"
  com.example.label-with-empty-value: ""
network: host
shm_size: '2gb'
target: prod
command: bundle exec thin -p 3000
configs:
  source: my_config
  target: /redis_config
  uid: '103'
  gid: '103'
  mode: 0440
container_name: custom_name
depends_on:
  db
  redis
deploy:
  mode: replicated
  replicas: 6
  endpoint_mode: vip
  labels:
    com.example.description: "This label will appear on the web service"
  constraints:
    node.role: manager
    engine.labels.operatingsystem: ubuntu 18.04
  preferences:
    spread: node.labels.zone
  placement:
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
```

```
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
  restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
  update_config:
    parallelism: 2
    delay: 10s
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
    - ./common.env
    - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost"]
    interval: 1m30s
    timeout: 10s
    retries: 3
    start_period: 40s
  logging:
    driver: "json-file"
    options:
      max-size: "200k"
      max-file: "10"
  networks:
    new:
      ipv4_address: 172.16.238.10
      ipv6_address: 2001:3984:3989::10
  ports:
```

```
- "172.16.238.10:5000:80"
```

```
profiles: frontend
```

```
redis:
```

```
  image: redis
```

```
  init: true
```

```
profiles: backend
```

```
db:
```

```
  image: postgres
```

```
  networks:
```

```
    new:
```

```
      aliases:
```

```
        - database
```

```
profiles: backend
```

```
configs:
```

```
  my_config:
```

```
    file: ./my_config.txt
```

```
  my_other_config:
```

```
    external: true
```

```
networks:
```

```
  new:
```

```
    ipam:
```

```
      driver: default
```

```
      config:
```

```
        - subnet: "172.16.238.0/24"
```

```
        - subnet: "2001:3984:3989::/64"
```

```
  legacy:
```

```
    ipam:
```

```
      driver: default
```

```
      config:
```

```
        - subnet: "172.16.238.25/48"
```

```
        - subnet: "2001:3984:3989::65/128"
```

```
profiles:
```

```
- frontend
```

```
- backend
```

```
- midtier
```


SECRETS

This is for secret/authentication/authorization details configuration.

We can either use short or long syntax for secrets section.

Long syntax has more options that will give more granularity

- Source
- Target
- Uid
- Mode

```
version: "3.9"

services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        alpine: latest
        corp/web_app: 3.14
    labels:
      com.example.description: "Accounting webapp"
      com.example.department: "Finance"
      com.example.label-with-empty-value: ""
    network: host
    shm_size: '2gb'
    target: prod
    command: bundle exec thin -p 3000
  configs:
    source: my_config
    target: /redis_config
```

```
    uid: '103'
    gid: '103'
    mode: 0440
  container_name: custom_name
  depends_on:
    db
    redis
  deploy:
    mode: replicated
    replicas: 6
    endpoint_mode: vip
    labels:
      com.example.description: "This label will appear on the web service"
    constraints:
      node.role: manager
      engine.labels.operatingsystem: ubuntu 18.04
    preferences:
      spread: node.labels.zone
    placement:
      max_replicas_per_node: 2
    update_config:
      parallelism: 2
      delay: 10s
    restart_policy:
      condition: on-failure
  resources:
    limits:
      cpus: '0.50'
      memory: 50M
    reservations:
      cpus: '0.25'
      memory: 20M
  restart_policy:
    condition: on-failure
    delay: 5s
    max_attempts: 3
    window: 120s
  update_config:
    parallelism: 2
    delay: 10s
```

```
    order: stop-first
  entrypoint: /code/entrypoint.sh
  env_file:
  - ./common.env
  - ./apps/web.env
  environment:
    RACK_ENV: development
    SHOW: true
  expose:
    - "3000"
  external_links:
    -project_db:db_mysql
  healthcheck:
    test: ["CMD", "curl", "-f", "http://localhost"]
    interval: 1m30s
    timeout: 10s
    retries: 3
    start_period: 40s
  logging:
    driver: "json-file"
    options:
      max-size: "200k"
      max-file: "10"
  networks:
    new:
      ipv4_address: 172.16.238.10
      ipv6_address: 2001:3984:3989::10
  ports:
    - "172.16.238.10:5000:80"
  profiles: frontend
  secrets:
    - source: my_secret
      target: redis_secret
    uid: '103'
    gid: '103'
    mode: 0440
  redis:
    image: redis
    init: true
  profiles: backend
```

```
db:
  image: postgres
  networks:
    new:
      aliases:
        - database
  profiles: backend
configs:
  my_config:
    file: ./my_config.txt
  my_other_config:
    external: true
networks:
  new:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.0/24"
        - subnet: "2001:3984:3989::/64"
  legacy:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.25/48"
        - subnet: "2001:3984:3989::65/128"
profiles:
  - frontend
  - backend
  - midtier
secrets:
  my_secret:
    file: ./my_secret.txt
  my_other_secret:
    external: true
```

VOLUMES

Mount container, build or named volumes with this option.
This is similar to `-mount,-v` option in docker run command

```
version: "3.9"
```

```
services:
  webapp:
    build:
      context: ./dir
      dockerfile: Dockerfile-alternate
      args:
        buildno: 1
      cache_from:
        alpine: latest
        corp/web_app: 3.14
      labels:
        com.example.description: "Accounting webapp"
        com.example.department: "Finance"
        com.example.label-with-empty-value: ""
      network: host
      shm_size: '2gb'
      target: prod
      command: bundle exec thin -p 3000
    configs:
      source: my_config
      target: /redis_config
      uid: '103'
      gid: '103'
      mode: 0440
    container_name: custom_name
    depends_on:
      db
      redis
    deploy:
      mode: replicated
      replicas: 6
      endpoint_mode: vip
      labels:
        com.example.description: "This label will appear on the web service"
      constraints:
        node.role: manager
        engine.labels.operatingsystem: ubuntu 18.04
      preferences:
        spread: node.labels.zone
      placement:
```

```
    max_replicas_per_node: 2
  update_config:
    parallelism: 2
    delay: 10s
  restart_policy:
    condition: on-failure
resources:
  limits:
    cpus: '0.50'
    memory: 50M
  reservations:
    cpus: '0.25'
    memory: 20M
restart_policy:
  condition: on-failure
  delay: 5s
  max_attempts: 3
  window: 120s
update_config:
  parallelism: 2
  delay: 10s
  order: stop-first
entrypoint: /code/entrypoint.sh
env_file:
- ./common.env
- ./apps/web.env
environment:
  RACK_ENV: development
  SHOW: true
expose:
- "3000"
external_links:
- project_db:db_mysql
healthcheck:
  test: ["CMD", "curl", "-f", "http://localhost"]
  interval: 1m30s
  timeout: 10s
  retries: 3
  start_period: 40s
logging:
```

```
driver: "json-file"
options:
  max-size: "200k"
  max-file: "10"
networks:
  new:
    ipv4_address: 172.16.238.10
    ipv6_address: 2001:3984:3989::10
ports:
  - "172.16.238.10:5000:80"
profiles: frontend
secrets:
  - source: my_secret
    target: redis_secret
    uid: '103'
    gid: '103'
    mode: 0440
volumes:
  - type: volume
    source: mydata
    target: /data
    volume:
      nocopy: true
  - type: bind
    source: ./static
    target: /opt/app/static
redis:
  image: redis
  init: true
  profiles: backend
db:
  image: postgres
  networks:
    new:
      aliases:
        - database
  profiles: backend
configs:
  my_config:
    file: ./my_config.txt
```

```
my_other_config:
  external: true
networks:
  new:
    ipam:
      driver: default
      attachable: true
      config:
        - subnet: "172.16.238.0/24"
        - subnet: "2001:3984:3989::/64"
  legacy:
    ipam:
      driver: default
      config:
        - subnet: "172.16.238.25/48"
        - subnet: "2001:3984:3989::65/128"
profiles:
  - frontend
  - backend
  - midtier
secrets:
  my_secret:
    file: ./my_secret.txt
  my_other_secret:
    external: true
volumes:
  mydata:
    driver_opts:
      type: "nfs"
      o: "addr=10.40.0.199,nolock,soft,rw"
      device: ":/docker/example"
  dbdata:
    external: true
```


DOCKER COMPOSE CLIENT COMMANDS

The terminal/cli commands that come in handy to run docker compose are as follows:

Commands:

build	Build or rebuild services
bundle	Generate a Docker bundle from the Compose file
config	Validate and view the Compose file
create	Create services
down	Stop and remove containers, networks, images, and volumes
events	Receive real time events from containers
exec	Execute a command in a running container
help	Get help on a command
images	List images
kill	Kill containers
logs	View output from containers
pause	Pause services
port	Print the public port for a port binding
ps	List containers
pull	Pull service images
push	Push service images
restart	Restart services
rm	Remove stopped containers
run	Run a one-off command
scale	Set number of containers for a service
start	Start services
stop	Stop services
top	Display the running processes

unpause	Unpause services
up	Create and start containers
version	Show the Docker-Compose version information

Docker diff:

This is used to view all the files that's changed in the read-write layer of a container

```
$ docker diff 1fd1f54c1b
```

Docker History

This is used to see how the image was built. It will have the details of all the layers as to how and when they were built

```
$ docker history docker
```

IMAGE	CREATED	CREATED BY	SIZE	COMMENT
3e23a5875458	8 days ago	/bin/sh -c #(nop) ENV LC_ALL=C.UTF-8	0 B	
8578938dd170	8 days ago	/bin/sh -c dpkg-reconfigure locales &&	loc	1.245 MB
be51b77efb42	8 days ago	/bin/sh -c apt-get update && apt-get install	338.3 MB	
4b137612be55	6 weeks ago	/bin/sh -c #(nop) ADD jessie.tar.xz in /	121 MB	
750d58736b4b	6 weeks ago	/bin/sh -c #(nop) MAINTAINER Tianon Gravi <ad	0 B	

Docker create

This helps to create a writable container layer on the specified image and prepares it for running a specific command

```
$ docker create -v /data --name data ubuntu
```

```
240633dfbb98128fa77473d3d9018f6123b99c454b3251427ae190a7d951ad57
```

```
$ docker run --rm --volumes-from data ubuntu ls -la /data
```

```
total 8
```

```
drwxr-xr-x  2 root root 4096 Dec  5 04:10 .
```

```
drwxr-xr-x 48 root root 4096 Dec  5 04:11 ..
```

Docker images

Lists all the images in the system.

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	77af4d6b9913	19 hours ago	1.089 GB
committ	latest	b6fa739cedf5	19 hours ago	1.089 GB
<none>	<none>	78a85c484f71	19 hours ago	1.089 GB
docker	latest	30557a29d5ab	20 hours ago	1.089 GB
<none>	<none>	5ed6274db6ce	24 hours ago	1.089 GB
postgres	9	746b819f315e	4 days ago	213.4 MB
postgres	9.3	746b819f315e	4 days ago	213.4 MB
postgres	9.3.5	746b819f315e	4 days ago	213.4 MB
postgres	latest	746b819f315e	4 days ago	213.4 MB

We can even list out specific images as :

```
$ docker images java:8
```

Docker inspect

This command helps us inspect almost anything. Be it network, image or container.

```
$ docker inspect c1
```

Where c1 is a container

Docker kill

This is used to kill one or more containers

```
$ docker kill my_container
```

Docker Network

This command is used to work with networks in docker daemon.

- Docker network create : used to create network

For instance:

```
$ docker network create -d overlay \  
--subnet=192.168.10.0/25 \  
--subnet=192.168.20.0/25 \  
--gateway=192.168.10.100 \  
--gateway=192.168.20.100 \  
--aux-address="my-router=192.168.10.5" --aux-address="my-switch=192.168.10.6" \  
--aux-address="my-printer=192.168.20.5" --aux-address="my-nas=192.168.20.6" \  
my-multihost-network
```

- Docker network connect: used to connect a container or multiple containers to a network

```
$ docker network connect --ip 10.10.36.122 \  
--link container1:c1 \  
--alias db --alias mysql \  
multi-host-network container2
```

- Docker network inspect: used to inspect the features/configuration of a network in docker daemon.
- Docker network ls: used to view all the networks available for us in the system.
- Docker network disconnect : Similar to connect, this option is used to disconnect a container or multiple containers from the network.
- Docker network prune: This will remove all the unused networks
- Docker network rm: This will remove all networks.

Docker node:

In case of orchestration, in docker swarm to be precise, this command is used to promote/demote/update a node.

Options available for us is :

- Docker node promote: promotes a node to become manager
- Docker node demote: demotes a node from being a manager
- Docker node ls: lists all nodes available
- Docker node ps: lists tasks running on one or more nodes
- Docker node update: updates a node
- Docker node rm: removes one or more nodes

Docker ps:

Used to list all the running containers. Along with option “-a”, it lists all the containers both running and exited.

```
$ Docker ps -a
```

Docker pull:

This command is used to pull the docker image from registry but explicitly isn't run to create a container of that image

```
$ docker pull debian
```

```
Using default tag: latest
```

```
latest: Pulling from library/debian
```

```
fdd5d7827f33: Pull complete
```

```
a3ed95caeb02: Pull complete
```

```
Digest: sha256:e7d38b3517548a1c71e41bffe9c8ae6d6d29546ce46bf62159837aad072c90aa
```

```
Status: Downloaded newer image for debian:latest
```

Docker push:

This is used to push an image to the docker registry.

```
$ docker image push --all-tags registry-host:5000/myname/myimage
```

```
The push refers to repository [registry-host:5000/myname/myimage]
```

```
195be5f8be1d: Pushed
```

```
latest: digest: sha256:edafc0a0fb057813850d1ba44014914ca02d671ae247107ca70c94db686e7de6 size: 4527
```

```
195be5f8be1d: Layer already exists
```

```
v1: digest: sha256:edafc0a0fb057813850d1ba44014914ca02d671ae247107ca70c94db686e7de6 size: 4527
```

```
195be5f8be1d: Layer already exists
```

```
v1.0: digest: sha256:edafc0a0fb057813850d1ba44014914ca02d671ae247107ca70c94db686e7de6 size: 4527
```

```
195be5f8be1d: Layer already exists
```

```
v1.0.1: digest: sha256:edafc0a0fb057813850d1ba44014914ca02d671ae247107ca70c94db686e7de6 size: 4527
```

Docker restart:

This command is used to restart one or more containers.

```
$ docker restart 195 fdd
```

Docker rm:

This command is used to remove the containers. If you want to forcefully remove a container, use “-f” along.

```
$ docker rm -v -f --link /webapp/redis
```

Docker rmi:

This is used to remove images from the Docker daemon

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
test1	latest	fd484f19954f	23 seconds ago	7 B (virtual 4.964 MB)
test	latest	fd484f19954f	23 seconds ago	7 B (virtual 4.964 MB)
test2	latest	fd484f19954f	23 seconds ago	7 B (virtual 4.964 MB)

```
$ docker rmi fd484f19954f
```

```
Error: Conflict, cannot delete image fd484f19954f because it is tagged in multiple repositories, use -f to force  
2013/12/11 05:47:16 Error: failed to remove one or more images
```

```
$ docker rmi test1:latest
```

```
Untagged: test1:latest
```

```
$ docker rmi test2:latest
```

```
Untagged: test2:latest
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
test	latest	fd484f19954f	23 seconds ago	7 B (virtual 4.964 MB)

```
$ docker rmi test:latest
```

```
Untagged: test:latest
```

```
Deleted: fd484f19954f4920da7ff372b5067f5b7ddb2fd3830cecd17b96ea9e286ba5b8
```

Docker Secret:

This command is used to create,inspect,remove secrets in the machine.

Create:

```
$ docker secret create \  
--label env=dev \  
--label rev=20170324 \  
my_secret ./secret.json  
eo7jnzguqgtpdah3cm5srfb97
```

Inspect:

```
$ docker secret inspect my_secret  
[  
  {  
    "ID": "eo7jnzguqgtpdah3cm5srfb97",  
    "Version": {  
      "Index": 17  
    },  
    "CreatedAt": "2017-03-24T08:15:09.735271783Z",  
    "UpdatedAt": "2017-03-24T08:15:09.735271783Z",  
    "Spec": {  
      "Name": "my_secret",  
      "Labels": {  
        "env": "dev",  
        "rev": "20170324"  
      }  
    }  
  }  
]
```

Is:

```
$ docker secret ls
```

ID	NAME	CREATED	UPDATED
6697bflskwj1998km1gnnjr38	q5s5570vtvnimefos1fyeo2u2	6 weeks ago	6 weeks ago

9u9hk4br2ej0wngngka6rp4hq	my_secret	5 weeks ago	5 weeks ago
mem02h8n73mybpgqjf0kfi1n0	test_secret	3 seconds ago	3 seconds ago

rm:

```
$ docker secret rm secret.json
saph4csdo5b6wz2p5uimh5xg
```

Docker service:

Docker service is used while orchestration. This is the equivalent to the service section in docker compose file.

All the options will be same of that of the docker compose file.

For instance:

```
$ docker service create \
  --name nginx \
  --replicas 2 \
  --replicas-max-per-node 1 \
  --placement-pref 'spread=node.labels.datacenter' \
  nginx
```

other options available are: create/inspect/ps/ls/rm/update.

Docker start/stop

These commands are used to start or stop a single or multiple containers.

```
$ docker start my_container
$ docker stop my_container
```

Docker swarm:

This is an orchestration tool used to manage multiple containers deployed across multiple host machines.

This is a world within docker. But for now we can go ahead with basic few commands :

Docker swarm ca

This command is used to view the current root CA certificate. In case if any node managers are compromised and no longer can be trusted, this certificate can be rotated with command “Docker swarm ca –rotate”

```
$ docker swarm ca --rotate
desired root digest: sha256:05da740cf2577a25224c53019e2cce99bcc5ba09664ad6bb2a9425d9ebd1b53e
rotated TLS certificates: [=====>] 2/2 nodes
rotated CA certificates: [=====>] 2/2 nodes
```


-----BEGIN CERTIFICATE-----

```
MIIBazCCARCgAwIBAgIUfYnG04h5Rr4IKyA4/E65tYKg8lwCgYIKoZlZj0EAwIw
EzERMA8GA1UEAxMlc3dhcm0tY2EwHhcNMTcwNTE2MDAxMDAwWhcNMzcwNTExMDAx
MDAwWjATMREwDwYDVQQDEwhzd2FybS1jYTBZMBMGByqGSM49AgEGCCqGSM49AwEH
A0IABC2DuNriETP7C7IfiEPk39tWaaU0I2RumUP4fX4+3m+87j0DU0CsemUaaOG6
+PxHhGu2VXQ4c9pctPHgf7vWeVajQjBAMA4GA1UdDwEB/wQEAwIBBjAPBgNVHRMB
Af8EBTADAQH/MB0GA1UdDgQWBBSSEL02z6mCI3SmMDmITMr12qCRY2jAKBggqhkJO
PQQDAgNJADBGAiEA263Eb52+825EeNQZM0AME+aoH1319Zp9/J5ijLW+6ACIQCg
gyg5u9Iliel99I7SuMhNeLkrU7fXs+Of1nTyyM73ig==
```

-----END CERTIFICATE-----

Along with detach, the progress of the rotation won't be displayed.

Docker swarm init:

This is used to initialize the swarm orchestration.

```
$ docker swarm init --advertise-addr 192.168.99.121
```

Swarm initialized: current node (bvz81updecsj6wjz393c09vti) is now a manager.

To add a worker to this swarm, run the following command:

```
docker swarm join \
--token SWMTKN-1-3pu6hszjas19xyp7ghgosyx9k8atbfc8p2is99znp26u2lkl-1awxwuwd3z9j1z3puu7rcgdbx \
172.17.0.2:2377
```

To add a manager to this swarm, run 'docker swarm join-token manager' and follow the instructions.

Docker swarm join-token

Join tokens are secrets that allow a node to the swarm.

```
$ docker swarm join-token worker
```

To add a worker to this swarm, run the following command:

```
docker swarm join \
--token SWMTKN-1-3pu6hszjas19xyp7ghgosyx9k8atbfc8p2is99znp26u2lkl-1awxwuwd3z9j1z3puu7rcgdbx \
172.17.0.2:2377
```

```
$ docker swarm join-token manager
```

To add a manager to this swarm, run the following command:

```
docker swarm join \
--token SWMTKN-1-3pu6hszjas19xyp7ghgosyx9k8atbfc8p2is99znp26u2lkl-7p73s1dx5in4tatdymyhg9hu2 \
```

```
172.17.0.2:2377
```

```
$ docker swarm join-token --rotate worker
```

Successfully rotated worker join token.

To add a worker to this swarm, run the following command:

```
$ docker swarm join \
```

```
--token SWMTKN-1-3pu6hszjas19xyp7ghgosyx9k8atbfc8p2is99znp26u2lkl-b30ljddcqh9b9v4rs7mel7t \
```

```
172.17.0.2:2377
```

Docker swarm join:

This is similar to join-token but without the secrets. If required, can be specified using --token.

```
$ docker swarm join --token SWMTKN-1-3pu6hszjas19xyp7ghgosyx9k8atbfc8p2is99znp26u2lkl-1awxwud3z9j1z3puu7rcgdbx 192.168.99.121:2377
```

This node joined a swarm as a worker.

```
$ docker node ls
```

ID	HOSTNAME	STATUS	AVAILABILITY	MANAGER STATUS
7ln70fl22uw2dvjn2ft53m3q5	worker2	Ready	Active	
dkp8vy1dq1kxleu9g4u78tlag	worker1	Ready	Active	Reachable
dvfxp4zseq4s0rih1selh0d20 *	manager1	Ready	Active	Leader

Docker swarm leave:

This is used to leave swarm.

```
$ docker swarm leave
```

Node left the default swarm.

Docker swarm unlock-key:

This command is used to unlock a swarm manager node once the docker daemon is restarted.

```
$ docker swarm unlock-key
```

To unlock a swarm manager after it restarts, run the `docker swarm unlock` command and provide the following key:

```
SWMKEY-1-fySn8TY4w5IKcWcJPkpKufejh9hxx5KYwx6XZigx3Q4
```

Please remember to store this key in a password manager, since without it you will not be able to restart the manager.

```
$ docker swarm unlock-key --rotate
```

Successfully rotated manager unlock key.

To unlock a swarm manager after it restarts, run the ``docker swarm unlock`` command and provide the following key:

```
SWMKEY-1-7c37Cc8654o6p38HnroywCi19pIIOnGtbdZEgtKxZu8
```

Please remember to store this key in a password manager, since without it you will not be able to restart the manager.

Docker swarm unlock:

This is used to unlock a manager using a user-supplied unlock keys.

```
$ docker swarm unlock
```

Please enter unlock key:

Docker swarm update:

This is used to update swarm with new parameters.

```
$ docker swarm update --cert-expiry 720h
```

Docker Volume

This command is used to create/inspect/prune/remove the volumes.

Create:

```
$ docker volume create --driver local \
  --opt type=tmpfs \
  --opt device=tmpfs \
  --opt o=size=100m,uid=1000 \
  Foo

$ docker volume create --driver local \
  --opt type=btrfs \
  --opt device=/dev/sda2 \
  Foo

docker volume create --driver local \
  --opt type=nfs \
  --opt o=addr=192.168.1.1,rw \
  --opt device=:/path/to/dir \
  foo
```

inspect:

used to inspect configurations of a volume in the system.

```
$ docker volume create myvolume
myvolume

$ docker volume inspect myvolume
[
  {
    "CreatedAt": "2020-04-19T11:00:21Z",
    "Driver": "local",
```

```
"Labels": {},
"Mountpoint":
"/var/lib/docker/volumes/8140a838303144125b4f54653b47ede0486282c623c3551fbc7f390cdc3e9cf5/_data",
"Name": "myvolume",
"Options": {},
"Scope": "local"
}
]

$ docker volume inspect --format '{{ .Mountpoint }}' myvolume
/var/lib/docker/volumes/myvolume/_data
```

ls:

Lists all the volumes

```
$ docker volume create rosemary
```

```
rosemary
```

```
$ docker volume create tyler
```

```
tyler
```

```
$ docker volume ls
```

DRIVER	VOLUME NAME
local	rosemary
local	tyler

prune:

removes unused volumes

```
$ docker volume prune
```

```
WARNING! This will remove all local volumes not used by at least one container.
```

```
Are you sure you want to continue? [y/N] y
```

```
Deleted Volumes:
```

```
07c7bdf3e34ab76d921894c2b834f073721fccfbcbba792aa7648e3a7a664c2e
```

```
my-named-vol
```

Total reclaimed space: 36 B

rm:

Removes a specific volume or volumes.

```
$ docker volume rm -f hello
```

hello

DOCKER ENGINE:

***Please note : Docker container PID mostly will be different on host to that of the container.

